

BEGINNER AND INTERMEDIATE LEVEL R COURSE

Marcos Jimenez

About the workshop

Materials can be accessed at https://github.com/Marcosjnez/R_USF.

1. R as a Calculator

R can be used as a calculator:

```
3 + 2          # Sum
```

```
[1] 5
```

```
5 - 2          # Subtraction
```

```
[1] 3
```

```
3 * 4          # Multiplication
```

```
[1] 12
```

```
15 / 5         # Division
```

```
[1] 3
```

```
2^3           # Power
```

```
[1] 8
```

```
10 + (2 + 3) * 3^2 # Preference order for computing
```

```
[1] 55
```

1.1. Assignments

R can store values using symbolic names. There are multiple ways to assign values:

```
a <- 3 # Best option  
3 -> a # Less common...  
a = 3  # Not recommended  
a      # Visualize the object
```

```
[1] 3
```

1.2. Reserved Words

Some names are reserved in R and cannot be used as variables:

```
# TRUE <- 4 # Error  
?reserved # See reserved words
```

```
starting httpd help server ... done
```

```
?objects
```

1.3. Symbolic Operations

You can use variables in arithmetic operations:

```
a <- 3  
b <- 2  
c <- a + b
```

2. Data Types in R

R supports several basic data types.

2.1. Numeric

```
a <- 1  
b <- 2  
c <- 1 + 2
```

2.2. Character Strings

```
d <- "cat"  
e <- "dog"
```

2.3. Factors

Factors are categorical variables:

```
d <- factor("cat")  
e <- "dog"
```

2.4. Logical Data

Logical operations return TRUE or FALSE:

```
A <- 5  
A == A # TRUE
```

```
[1] TRUE
```

```
B <- 5  
A == B # TRUE
```

```
[1] TRUE
```

```
a <- 4  
A == a # FALSE: R is case-sensitive
```

```
[1] FALSE
```

```
A > a # TRUE: A is greater than a
```

```
[1] TRUE
```

2.5. Complex Numbers

R supports complex numbers:

```
i <- 1i          # Square root of -1  
exp(i*pi) + 1    # Euler's identity
```

```
[1] 0+1.224606e-16i
```

2.6. Constants

R has built-in constants like pi:

```
?Constants
```

```
pi
```

```
[1] 3.141593
```

2.1. Vectors

Vectors can hold elements of the same type:

```
n <- c(1, 2, 3, 4)           # Numeric
l <- c(TRUE, FALSE)          # Logical
k <- c("cat", "dog")          # Character
k <- factor(c("cat", "dog"), ordered = TRUE) # Factor
x <- c(1 + 2i, 4i)            # Complex

# Preference order:
# 1. Character
# 2. Complex
# 3. Numeric
# 4. Logical / Factor
```

2.2. Constructing Vectors

```
# Vector of numbers from 0 to 20 by steps of 2:
y <- seq(from = 0, to = 20, by = 2)
w <- 1:4 # Vector of integers from 1 to 4

# Concatenate the vectors:
z <- c(y, w)

# Take samples from random variables:
set.seed(2025)
rbinom(n = 10, size = 5, prob = 0.5)
```

```
[1] 3 2 3 2 3 3 4 1 3 3
```

```
rnorm(n = 10, mean = 0, sd = 1)
```

```
[1] -0.16285434  0.39711189 -0.07998932 -0.34496518  0.70215136 -0.39569639  
[7] -1.75505405 -0.42096376  0.76490961  1.06616211
```

```
rt(n = 10, df = 5)
```

```
[1] -0.2206268 -2.0056293 -0.2171123  2.5775456 -0.2399014 -1.4348883  
[7]  1.2676429 -0.6937007 -0.7065318  0.6663071
```

```
rchisq(n = 10, df = 5)
```

```
[1] 4.792938 4.666413 2.102163 2.458308 1.923977 7.847497 2.850527 5.289288  
[9] 6.635302 2.455811
```

```
runif(n = 10, min = 0, max = 1)
```

```
[1] 0.1097217 0.8087052 0.4355751 0.8042243 0.8617182 0.4337776 0.6402778  
[8] 0.9299222 0.8570572 0.1363258
```

```
# Random letters and categories:
```

```
sample(x = letters, size = 10, replace = TRUE)
```

```
[1] "q" "o" "m" "d" "n" "n" "l" "t" "w" "i"
```

```
sample(c("dog", "cat"), size = 5, replace = TRUE)
```

```
[1] "dog" "dog" "dog" "cat" "dog"
```

```
# Automatic creation of long character vectors:
```

```
paste("Sujeto", 1:20, sep = "")
```

```
[1] "Sujeto1" "Sujeto2" "Sujeto3" "Sujeto4" "Sujeto5" "Sujeto6"  
[7] "Sujeto7" "Sujeto8" "Sujeto9" "Sujeto10" "Sujeto11" "Sujeto12"  
[13] "Sujeto13" "Sujeto14" "Sujeto15" "Sujeto16" "Sujeto17" "Sujeto18"  
[19] "Sujeto19" "Sujeto20"
```

2.3. Vector Indexing

```
z <- 1:6
```

```
# Extract elements:
```

```
z[1]
```

```
[1] 1
```

```
z[c(2, 5)]
```

```
[1] 2 5
```

```
# Remove elements:
```

```
z[-1]
```

```
[1] 2 3 4 5 6
```

```
z[-c(2, 5)]
```

```
[1] 1 3 4 6
```

2.4. Numeric Vector Functions

```
x <- sample(x = 1:5, size = 6, replace = TRUE)
```

```
mean(x)
```

```
[1] 3.333333
```

```
median(x)
```

```
[1] 4
```

```
var(x)
```

```
[1] 2.266667
```

```
sd(x)
```

```
[1] 1.505545
```

```
min(x)
```

```
[1] 1
```

```
max(x)
```

```
[1] 5
```

```
sum(x)
```

```
[1] 20
```

```
prod(x)
```

```
[1] 640
```

```
log(x)
```

```
[1] 1.3862944 1.3862944 0.6931472 1.6094379 0.0000000 1.3862944
```

```
exp(x)
```

```
[1] 54.598150 54.598150 7.389056 148.413159 2.718282 54.598150
```

```
length(x)
```

```
[1] 6
```



```
head(x, n = 3)
```

```
[1] 4 4 2
```

```
tail(x, n = 3)
```

```
[1] 5 1 4
```

```
table(x)
```


```
x  
1 2 4 5  
1 1 3 1
```

```
?mean  
x[1] <- NA  
mean(x, na.rm = TRUE)
```

```
[1] 3.2
```

2.4.1. Plotting Two Numeric Vectors

```
x <- rnorm(n = 100, mean = 100, sd = 15)  
y <- rnorm(n = 100, mean = 100, sd = 15)  
  
plot(x, y, ylab = "Coordinate 1", xlab = "Coordinate 2")  
  
# Highlight specific coordinates:  
condition <- x > 100 & x < 115 & y > 100 & y < 115  
indices <- which(condition)  
points(x[indices], y[indices], col = "red")  
  
# Draw segments:  
segments(x0 = 100, x1 = 100, y0 = 0, y1 = 150, lty = "dashed")  
segments(x0 = 115, x1 = 115, y0 = 0, y1 = 150, lty = "dashed")  
segments(x0 = 0, x1 = 150, y0 = 100, y1 = 100, lty = "dashed")  
segments(x0 = 0, x1 = 150, y0 = 115, y1 = 115, lty = "dashed")
```



2.5. Simulating Random Variables

2.5.1. Binomial Distribution


```
set.seed(123)                                # Random seed
N <- 1000                                     # Sample size
n <- 8                                        # Number of trials
prob <- 0.60                                 # Probability of success
b <- rbinom(N, size = n, prob = prob)        # 1000 samples from Binomial
empirical <- table(b)/1000                   # Empirical proportions

factual <- dbinom(0:n, size = n,             # Theoretical proportions
                  prob = prob)
cbind(empirical, factual)                   # Compare empirical and theoretical
```

	empirical	factual
0	0.001	0.00065536
1	0.007	0.00786432
2	0.044	0.04128768
3	0.112	0.12386304
4	0.235	0.23224320
5	0.287	0.27869184
6	0.206	0.20901888

```
7      0.092 0.08957952
8      0.016 0.01679616
```

```
# Plot the theoretical distribution:
plot(0:n, factorial, main = paste("X ~ B(n=", n, ", p=", prob, ")"), sep = ""),
     xlab = "X", ylab = "Probability", type = "h")
```



2.5.2. Normal Distribution

```
X <- rnorm(n = 100000, mean = 0, sd = 1)

x <- 1.25
delta <- 1
mean((X > x-delta/2 & X < x+delta/2)) / delta
```

```
[1] 0.18677
```

```
dnorm(x, mean = 0, sd = 1)
```

```
[1] 0.1826491
```


```

curve(dnorm(x, mean = 0, sd = 1), xlim = c(-3, 3), ylab = "Density")
segments(x0 = x-delta/2, x1 = x-delta/2, y0 = 0, y1 = dnorm(x-delta/2))
segments(x0 = x+delta/2, x1 = x+delta/2, y0 = 0, y1 = dnorm(x+delta/2))

x_fill <- seq(x - delta/2, x + delta/2, length.out = 100)
y_fill <- dnorm(x_fill, mean = 0, sd = 1)

polygon(c(x - delta/2, x_fill, x + delta/2),
        c(0, y_fill, 0),
        col = "salmon", border = NA)

```



2.5.3. Chi-Square Distribution

```

df <- 5
X <- rchisq(n = 100000, df = df)

x <- 8
delta <- 1
mean((X > x-delta/2 & X < x+delta/2)) / delta

```

```
[1] 0.05497
```


```
dchisq(x, df = df)
```

```
[1] 0.05511196
```

```
curve(dchisq(x, df = df), xlim = c(0, 15), ylab = "Density")
segments(x0 = x-delta/2, x1 = x-delta/2, y0 = 0, y1 = dchisq(x-delta/2, df = df))
segments(x0 = x+delta/2, x1 = x+delta/2, y0 = 0, y1 = dchisq(x+delta/2, df = df))

x_fill <- seq(x - delta/2, x + delta/2, length.out = 100)
y_fill <- dchisq(x_fill, df = df)

polygon(c(x - delta/2, x_fill, x + delta/2),
        c(0, y_fill, 0),
        col = "salmon", border = NA)
```



2.5.4. From normal to Chi2 to t-student distribution

```
nsim <- 1000 # Number of simulations
X2 <- vector(length = nsim) # Store the Chi2 values
t <- vector(length = nsim) # Store the t values
df <- 9 # Degrees of freedom
```

```

for(i in 1:nsim) {


  z <- rnorm(df)           # Sample from standard normal
  X2[i] <- sum(z^2)         # Sum the squares of the normals (Chi2)
  Z <- rnorm(1)            # Take one sample from the standard normal
  t[i] <- Z / sqrt(X2[i]/df) # Build a variable that is t-distributed

}

hist(X2, freq = FALSE) # Plot Chi2 distribution
curve(dchisq(x, df = df), add = TRUE)

```


Histogram of X2



```

hist(t, freq = FALSE) # Plot t-student distribution
curve(dt(x, df = df), add = TRUE)

```



3.1. Creation of Matrices

```
# Create a matrix with a number of rows and columns:
N <- matrix(1:4, nrow = 2, ncol = 2)
N
```

```
      [,1] [,2]
[1,]    1    3
[2,]    2    4
```

```
# Create a matrix by binding vectors:
x1 <- rnorm(n = 5, mean = 0, sd = 1)
x2 <- runif(n = 5, min = 0, max = 1)
rbind(x1, x2)
```

```
      [,1]      [,2]      [,3]      [,4]      [,5]
x1 0.1634255 -0.7646141 -1.0209128 -0.4881263  1.394205
x2 0.9057754  0.5655829  0.7639358  0.5377105  0.290256
```

```
X <- cbind(x1, x2)
X
```

```
      x1      x2
[1,] 0.1634255 0.9057754
[2,] -0.7646141 0.5655829
[3,] -1.0209128 0.7639358
[4,] -0.4881263 0.5377105
[5,] 1.3942046 0.2902560
```

```
# Set names:
colnames(X) <- paste("Score", 1:2, sep = "_")
rownames(X) <- paste("Subject", 1:5, sep = " ")
X
```

```
      Score_1 Score_2
Subject 1 0.1634255 0.9057754
Subject 2 -0.7646141 0.5655829
Subject 3 -1.0209128 0.7639358
Subject 4 -0.4881263 0.5377105
Subject 5 1.3942046 0.2902560
```

3.2. Matrix Indexing

```
X <- matrix(rnorm(2*4), nrow = 2, ncol = 4)
X
```

```
      [,1]      [,2]      [,3]      [,4]
[1,] 0.08442272 0.8206262 0.63070660 -1.0677470
[2,] 0.64935970 0.1645163 -0.05723784 0.4013795
```

```
X[2, 4] # Element in row 2, column 4
```

```
[1] 0.4013795
```



```
X[2, ] # All elements in row 2
```

```
[1] 0.64935970 0.16451629 -0.05723784 0.40137953
```

```
X[, 4] # All elements in column 4
```

```
[1] -1.0677470 0.4013795
```

3.3. Matrix Operations

```
# Matrix by vector multiplication:  
set.seed(123)  
N <- 20  
p <- 4  
b <- rnorm(n = p, mean = 0.5, sd = 2)  
X <- matrix(rnorm(N*p), nrow = N, ncol = p-1)
```

Warning in matrix(rnorm(N * p), nrow = N, ncol = p - 1): data length [80] is not a sub-multiple or multiple of the number of columns [3]

```
X <- cbind(1, X)  
  
e <- rnorm(N, mean = 0, sd = 1.5)  
y <- X %*% b + e  
  
fit <- lm(y ~ 0 + X)  
summary(fit)
```

Call:

```
lm(formula = y ~ 0 + X)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.3744	-0.5936	-0.2121	0.7138	2.3042

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
X1	-0.19338	0.27095	-0.714	0.486
X2	0.02804	0.28091	0.100	0.922
X3	3.91972	0.29229	13.410	4.05e-10 ***
X4	-0.10644	0.34319	-0.310	0.760

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.201 on 16 degrees of freedom

Multiple R-squared: 0.9193, Adjusted R-squared: 0.8991

F-statistic: 45.58 on 4 and 16 DF, p-value: 1.5e-08

```
fit$coefficients
```

	X1	X2	X3	X4
	-0.19338130	0.02804353	3.91971708	-0.10644232

```
solve(t(X) %*% X) %*% t(X) %*% y
```

	[,1]
[1,]	-0.19338130
[2,]	0.02804353
[3,]	3.91971708
[4,]	-0.10644232

3.3. Matrix by Matrix Multiplication

```
Z <- matrix(rnorm(6), nrow = 3, ncol = 2)
Y <- matrix(rnorm(6), nrow = 2, ncol = 3)
ZY <- Z %*% Y
ZY
```

	[,1]	[,2]	[,3]
[1,]	-0.4665383	1.63228136	-0.9965839
[2,]	-0.2052576	0.09397571	-0.1378942
[3,]	1.0103096	1.21882211	-0.1309261

```
Z[, 1, drop = FALSE] %*% Y[1, , drop = FALSE] +
  Z[, 2, drop = FALSE] %*% Y[2, , drop = FALSE]
```

```
      [,1]      [,2]      [,3]
[1,] -0.4665383  1.63228136 -0.9965839
[2,] -0.2052576  0.09397571 -0.1378942
[3,]  1.0103096  1.21882211 -0.1309261
```

3.3. Matrix Utilities

```
ncol(ZY)
```

```
[1] 3
```

```
nrow(ZY)
```

```
[1] 3
```

```
rowSums(ZY)
```

```
[1]  0.1691592 -0.2491761  2.0982055
```

```
colMeans(ZY)
```

```
[1]  0.1128379  0.9816931 -0.4218014
```

```
# Apply a custom function to each column
X <- matrix(runif(16), nrow = 4, ncol = 4)

geomean <- function(x) {
  n <- length(x)
  result <- prod(x)^(1/n)
  return(result)
}

apply(X, MARGIN = 2, FUN = geomean)
```

```
[1] 0.2393893 0.3770300 0.4454315 0.5436993
```

3.4. Arrays

```
A <- array(rnorm(2*3*4), dim = c(2, 3, 4))  
A
```

```
, , 1
```

	[,1]	[,2]	[,3]
[1,]	1.8438620	0.23538657	-0.96185663
[2,]	-0.6519499	0.07796085	-0.07130809

```
, , 2
```

	[,1]	[,2]	[,3]
[1,]	1.4445509	0.04123292	-2.053247
[2,]	0.4515041	-0.42249683	1.131337

```
, , 3
```

	[,1]	[,2]	[,3]
[1,]	-1.4606401	1.909104	0.7017843
[2,]	0.7399475	-1.443893	-0.2621975

```
, , 4
```

	[,1]	[,2]	[,3]
[1,]	-1.572144	-1.6015362	-1.4617556
[2,]	-1.514668	-0.5309065	0.6879168

```
apply(A, MARGIN = 2, FUN = mean)      # Across dimensions 1 and 3
```

```
[1] -0.08994217 -0.21689360 -0.28616584
```

```
apply(A, MARGIN = c(1, 2), FUN = sum) # Across slices
```

```
      [,1]      [,2]      [,3]
[1,] 0.2556286 0.5841869 -3.775075
[2,] -0.9751660 -2.3193357  1.485748
```

3.5. Data Frames

```
N <- 100
x1 <- sample(c("dog", "cat"), size = N, replace = TRUE)
x2 <- sample(1:4, size = N, replace = TRUE)

df <- data.frame(pet = x1, score = x2)
df$pet
```

```
[1] "cat" "cat" "dog" "dog" "cat" "cat" "dog" "cat" "dog" "dog" "cat" "dog"
[13] "dog" "dog" "cat" "dog" "cat" "dog" "cat" "cat" "cat" "cat" "dog" "dog"
[25] "dog" "cat" "dog" "dog" "cat" "dog" "dog" "cat" "dog" "cat" "cat" "cat"
[37] "dog" "dog" "dog" "dog" "cat" "dog" "dog" "dog" "dog" "cat" "cat" "dog"
[49] "cat" "dog" "dog" "dog" "cat" "dog" "dog" "dog" "cat" "dog" "dog" "cat"
[61] "dog" "dog" "dog" "cat" "cat" "dog" "dog" "dog" "cat" "dog" "cat" "dog"
[73] "cat" "dog" "dog" "dog" "cat" "cat" "dog" "dog" "cat" "dog" "dog" "cat"
[85] "dog" "cat" "cat" "cat" "cat" "cat" "cat" "cat" "cat" "cat" "dog" "cat"
[97] "dog" "dog" "dog" "dog"
```

```
df$score
```

```
[1] 3 3 4 2 3 3 3 1 1 3 2 3 2 1 4 2 2 2 4 1 4 1 1 3 3 1 1 3 1 3 4 2 1 1 3 4 2
[38] 3 2 1 2 3 2 3 3 1 2 4 1 2 2 4 2 4 2 1 4 2 4 1 4 2 1 4 4 3 2 1 3 2 4 3 2 3
[75] 3 1 4 2 3 1 1 3 2 3 2 2 2 4 1 4 2 1 4 4 1 4 1 4 2 2
```

```
fit <- lm(score ~ pet, data = df)
summary(fit)
```

```

Call:
lm(formula = score ~ pet, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-1.4386 -0.6836 -0.4186  0.5814  1.5814

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  2.41860     0.16769   14.42  <2e-16 ***
petdog       0.01999     0.22211    0.09   0.928
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.1 on 98 degrees of freedom
Multiple R-squared:  8.266e-05, Adjusted R-squared:  -0.01012
F-statistic: 0.008101 on 1 and 98 DF,  p-value: 0.9285

```

3.6. Lists

```

L <- list(1, 1:5, matrix(1:16, nrow = 4, ncol = 4))
L

```

```

[[1]]
[1] 1

[[2]]
[1] 1 2 3 4 5

[[3]]
      [,1] [,2] [,3] [,4]
[1,]    1    5    9   13
[2,]    2    6   10   14
[3,]    3    7   11   15
[4,]    4    8   12   16

```

```
# Multiply and sum rank-1 matrices using lapply
f <- function(i) {
  x <- Z[, i, drop = FALSE] %*% Y[i, , drop = FALSE]
  return(x)
}
ZY <- lapply(1:ncol(Z), FUN = f)
ZY
```

```
[[1]]
      [,1]      [,2]      [,3]
[1,] 0.54751085 1.53960722 -0.49427754
[2,] 0.02590656 0.07284958 -0.02338772
[3,] 0.45159240 1.26988337 -0.40768504
```

```
[[2]]
      [,1]      [,2]      [,3]
[1,] -1.0140492 0.09267413 -0.5023063
[2,] -0.2311642 0.02112613 -0.1145065
[3,] 0.5587172 -0.05106126 0.2767589
```

```
ZY[[1]] + ZY[[2]]
```

```
      [,1]      [,2]      [,3]
[1,] -0.4665383 1.63228136 -0.9965839
[2,] -0.2052576 0.09397571 -0.1378942
[3,] 1.0103096 1.21882211 -0.1309261
```

3.7. Objects in the Environment

```
objects()
```

```
[1] "a"      "A"      "b"      "B"      "c"      "condition"
[7] "d"      "delta"  "df"     "e"      "empirical" "f"
[13] "factual" "fit"    "geomean" "i"      "indices"  "k"
[19] "l"      "L"      "n"      "N"      "nsim"     "p"
[25] "prob"   "t"      "w"      "x"      "X"        "x_fill"
[31] "x1"     "x2"     "X2"     "y"      "Y"        "y_fill"
[37] "z"      "Z"      "ZY"
```

```
ls()
```

```
[1] "a"      "A"      "b"      "B"      "c"      "condition"
[7] "d"      "delta"  "df"     "e"      "empirical" "f"
[13] "factual" "fit"    "geomean" "i"      "indices"  "k"
[19] "l"      "L"      "n"      "N"      "nsim"     "p"
[25] "prob"   "t"      "w"      "x"      "X"        "x_fill"
[31] "x1"     "x2"     "X2"     "y"      "Y"        "y_fill"
[37] "z"      "Z"      "ZY"
```

4.1. Simulation Setup for the Linear Model

```
set.seed(123) # Fix the seed to get the same random numbers

# Setup:
N <- 10 # Sample size
x <- rnorm(N, mean = 0, sd = 1) # Predictor / Independent variable
sigma <- 1.5 # Standard deviation of the errors
true_se <- sigma / sqrt(var(x)*(N-1)) # True standard error of b
power <- 0.80 # Statistical power
alpha <- 0.05 # Type-I error
threshold <- qnorm(1-alpha) # One-sided point for rejection

# Get the b that gives the desired statistical power:
b <- qnorm(power, mean = threshold, sd = 1) * true_se
```

4.2. Simulating Power

```
# Simulation setup:
nsim <- 1000 # Number of replicas
pval <- vector(length = nsim) # Initialize the vector

for(i in 1:nsim) {
```



```

e <- rnorm(N, mean = 0, sd = sigma)
y <- 1 + x*b + e
fit <- lm(y ~ x)
z_statistic <- fit$coefficients[2] / true_se
pval[i] <- 1-pnorm(z_statistic, mean = 0, sd = 1)
}

mean(pval < alpha) # Proportion of times that the pvalue is smaller than alpha

```

```
[1] 0.79
```

```
power
```

```
[1] 0.8
```

4.3. Visual Insight into Statistical Power

```

z_statistic <- b / true_se

# Visualize the null hypothesis:
curve(dnorm(x, mean = 0, sd = 1), xlim = c(-4, 6), lwd = 2,
      xlab = "Z", ylab = "Density", main = "Null vs. Alternative hypothesis")

# Visualize the alternative hypothesis:
curve(dnorm(x, mean = z_statistic, sd = 1), col = "salmon", lwd = 2, add = TRUE)

# Shaded area under the alternative curve beyond the threshold:
x_fill <- seq(threshold, 10, length.out = 100)
y_fill <- dnorm(x_fill, mean = z_statistic, sd = 1)

polygon(c(threshold, x_fill, 10),
        c(0, y_fill, 0),
        col = "salmon", border = NA)

# Highlight the true z-statistic:

```


```

segments(x0 = z_statistic,
         x1 = z_statistic,
         y0 = 0,
         y1 = dnorm(z_statistic, mean = z_statistic),
         lty = "dashed")

axis(1, at = z_statistic, label = "z")
legend(x = "topleft", legend = c(paste("alpha =", alpha),
                                paste("power =", power)))

```

Null vs. Alternative hypothesis



4.4. Visual Insight into the P-value

```

# Generate a random z-statistic under the alternative hypothesis:
z_statistic <- rnorm(1, mean = b / true_se, sd = 1)
pval <- round(1 - pnorm(z_statistic), 3)

curve(dnorm(x, mean = 0, sd = 1), xlim = c(-4, 4), lwd = 2,
      xlab = "Z", ylab = "Density", main = "Null hypothesis")

# Shade the p-value region:
x_fill <- seq(z_statistic, 10, length.out = 100)

```

```

y_fill <- dnorm(x_fill, mean = 0, sd = 1)


polygon(c(z_statistic, x_fill, 10),
       c(0, y_fill, 0),
       col = "skyblue", border = NA)

segments(x0 = threshold,
        x1 = threshold,
        y0 = 0,
        y1 = dnorm(threshold, mean = 0, sd = 1),
        lty = "dashed")

axis(1, at = z_statistic, label = "z")
legend(x = "topleft", legend = c(paste("alpha =", alpha),
                                paste("p-value =", pval)))

```

Null hypothesis



5.1. Linear Model with Multiple Predictors

```

set.seed(123)
N <- 20
p <- 4

```

```
b <- rnorm(n = p, mean = 0.5, sd = 2)
X <- matrix(rnorm(N*p), nrow = N, ncol = p-1)
```

Warning in matrix(rnorm(N * p), nrow = N, ncol = p - 1): data length [80] is not a sub-multiple or multiple of the number of columns [3]

```
X <- cbind(1, X)

nsim <- 1000
sigma <- 1.5
true_se <- sqrt(diag(sigma^2 * solve(t(X) %*% X)))
power <- c(0.80, 0.60, 0.40, 0.05)
alpha <- 0.05
threshold <- qt(1 - alpha, df = N - p)

f <- function(power, threshold) {
  suppressWarnings(uniroot(\(x) (1 - pt(threshold, df = N - p, ncp = x)) - power,
    interval = c(-6, 6))$root)
}

t_stat <- sapply(power, FUN = f, threshold = threshold)
b <- t_stat * true_se

# Initialize result storage:
coefs <- matrix(NA, nsim, p)
se <- matrix(NA, nsim, p)
ts <- matrix(NA, nsim, p)
pval <- matrix(NA, nsim, p)
upper <- matrix(NA, nsim, p)
lower <- matrix(NA, nsim, p)
error <- vector(length = nsim)

for (i in 1:nsim) {
  e <- rnorm(N, mean = 0, sd = sigma)
  y <- X %*% b + e
  fit <- lm(y ~ 0 + X)
  coefs[i, ] <- coefficients(fit)
  error[i] <- sum(resid(fit)^2) / (N - p)
  se[i, ] <- sqrt(diag(error[i] * solve(t(X) %*% X)))
  ts[i, ] <- coefs[i, ] / se[i, ]
  upper[i, ] <- coefficients(fit) + qnorm(1 - alpha/2) * se[i, ]
  lower[i, ] <- coefficients(fit) + qnorm(alpha/2) * se[i, ]
}
```

```
pval[i, ] <- (1 - pt(ts[i, ], df = N - p))
}
```

5.1. Checking Simulation Accuracy

```
cbind(colMeans(coefs), b)      # Linear coefficients
```

```

              b
[1,] 0.87159473 8.793200e-01
[2,] 0.69604900 6.956320e-01
[3,] 0.52666068 5.304025e-01
[4,] -0.02259479 -7.209672e-08
```

```
cbind(colMeans(se), true_se)  # Standard errors
```

```

          true_se
[1,] 0.3328807 0.3383593
[2,] 0.3451118 0.3507917
[3,] 0.3590934 0.3650034
[4,] 0.4216360 0.4285753
```

```
cbind(colMeans(ts), t_stat)   # t statistics
```

```

          t_stat
[1,] 2.69519451 2.598776e+00
[2,] 2.08044868 1.983035e+00
[3,] 1.51452823 1.453144e+00
[4,] -0.05401576 -1.682241e-07
```

```
apply(pval, MARGIN = 2, FUN = \(x) mean(x < alpha)) # Empirical power
```

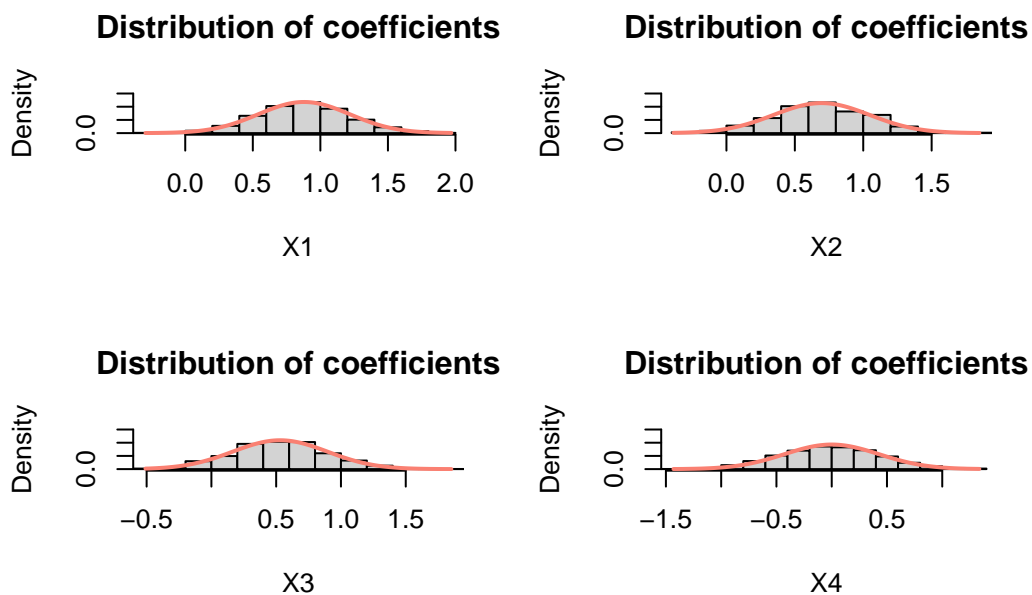
```
[1] 0.799 0.585 0.398 0.045
```

```
power
```

```
[1] 0.80 0.60 0.40 0.05
```

5.2. Distribution of Coefficients

```
par(mfrow = c(2, 2))
for (i in 1:p) {
  hist(coefs[, i], main = "Distribution of coefficients", freq = FALSE,
       xlim = range(coefs[, i]), ylim = c(0, 1.5),
       xlab = names(coefficients(fit))[i])
  curve(dnorm(x, b[i], sd = true_se[i]), lwd = 2, col = "salmon", add = TRUE)
}
```

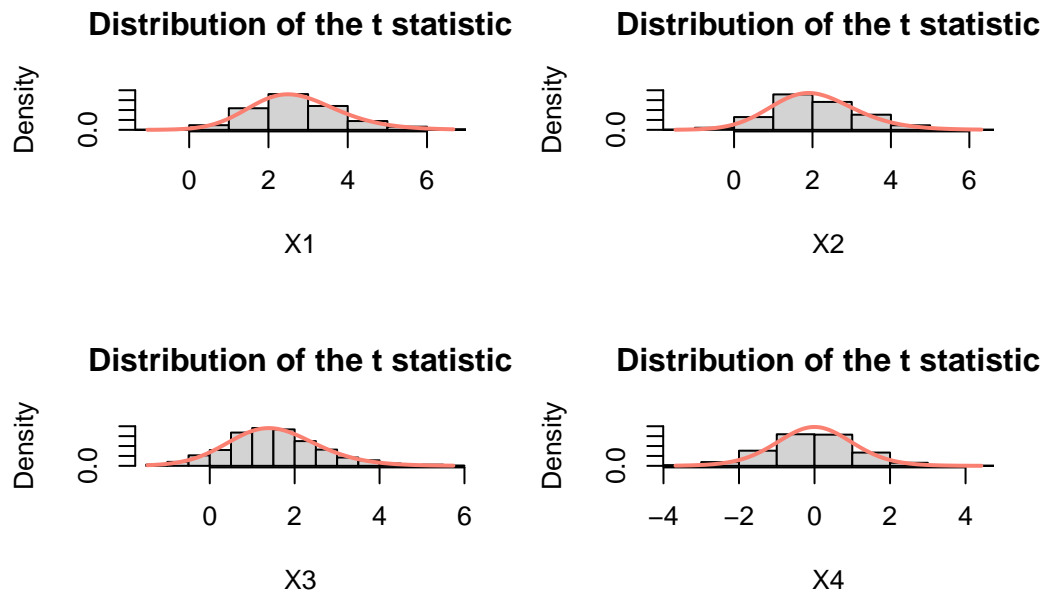


5.3. Distribution of t Statistics

```

par(mfrow = c(2, 2))
for (i in 1:p) {
  hist(ts[, i], main = "Distribution of the t statistic", freq = FALSE,
       xlim = range(ts[, i]), ylim = c(0, 0.4),
       xlab = names(coefficients(fit))[i])
  curve(dt(x, df = N - p, ncp = t_stat[i]), lwd = 2, col = "salmon", add = TRUE)
}

```




5.4. P-curve

```


par(mfrow = c(2, 2))
for (i in 1:p) {
  title <- paste("Distribution of the p-value for power =", power[i])
  hist(pval[, i], main = title, freq = FALSE,
       xlim = c(0, 1), xlab = names(coefficients(fit))[i])
}

```

istribution of the p-value for poweristribution of the p-value for power



istribution of the p-value for poweristribution of the p-value for power



5.5. The Dance of Confidence Intervals

```
par(mfrow = c(1, 1))
success <- rep(0, p)
failure <- rep(0, p)
nsim <- 20

for (i in 1:nsim) {
  plot(b, 1:p, xlab = "Value", ylab = "Coefficient",
       ylim = c(-1, 4.5), xlim = c(-2, 5),
       main = "The Dance of confidence intervals",
       yaxt = "n")
  axis(2, labels = paste(1:p), at = 1:p)

  for (j in 1:p) {
    segments(x0 = upper[i, j], x1 = lower[i, j], y0 = j, y1 = j)

    if (b[j] > lower[i, j] & b[j] < upper[i, j]) {
      points(b[j], j, bg = "skyblue", pch = 21)
      success[j] <- success[j] + 1
    } else {
```



```
    points(b[j], j, bg = "red", pch = 21)
    failure[j] <- failure[j] + 1
  }

  text(x = 4.5, y = j, labels = paste(success[j], failure[j], sep = "/"))
}

text(x = 1, y = 0, paste("Iteration", i, sep = "="))
Sys.sleep(2)
}
```